

Conference Paper Title

1st Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

2nd Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

3rd Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

4th Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

5th Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

5th Given Name Surname

dept. name of organization (of Aff.)

name of organization (of Aff.)

City, Country

email address or ORCID

Abstract—Citation networks are graphical networks composed of publications as nodes and citations as edges. They offer a unique look into scientific collaboration and the flow of knowledge within academia. In this paper, we examine the history of machine learning hardware papers through the lens of citation networks and construct GNN models for node classification and link prediction for our custom machine learning hardware dataset as well as existing datasets in software engineering and machine learning. We utilize existing citation graphs as well as curate our own dataset from the Web of Science database to show that node classification and link prediction are effective on a new network of machine learning hardware papers.

I. INTRODUCTION AND MOTIVATION

Inspired by the recent advent of graphical neural networks (GNNs), this project explores their potential to classify academic networks. Given any individual publication or author, GNNs could predict their relation to the remainder of the network, useful for understanding the relationships between academic fields, authors, institutions, and more. This opens the door to opportunities for classifying papers in fields with minimal representation in existing citation networks or predicting links between papers, potentially applicable to discovering inspiring papers related to an author’s past work, among other capabilities.

In particular, machine learning hardware is a new field that has skyrocketed from the recent popularity of machine learning. Due to the compute-intensive nature of this field, plenty of research exists in the hardware infrastructure to support these computing loads. We utilize GNNs for node classification and link prediction to explore existing datasets and comment on the new corpus of machine learning hardware papers. The main challenges are processing data for useful representations of the network and data interpretation on the created GNNs to comment on the field of machine learning hardware.

Identify applicable funding agency here. If none, delete this.

II. PREVIOUS WORK

There have been many previous explorations of citation networks using the datasets Cora, Citeseer, and PubMed Diabetes. In these citation networks, in particular, node classification and link prediction have been heavily explored. In Optimization of Graph Neural Networks with Natural Gradient Descent, Izadi et al improve on traditional optimization algorithms such as ADAM and stochastic gradient descent (SGD), producing a superior accuracy ($90.16 \pm 0.59\%$) by utilizing a proposed method inspired by natural gradient descent [1]. This paper achieved the highest accuracy across 71 different papers done on node classification of the Cora dataset.

Those that explored link prediction in the Cora and Citeseer datasets have achieved great results as well. In ‘NESS: Node Embeddings from Static SubGraphs,’ Ucar proposed splitting up the graph into multiple subgraphs that do not have overlapping edges in between subgraphs, using each graph to train and obtain node representations, then aggregating the results to obtain predictions [2]. Using this, he was able to obtain state-of-the-art accuracy results, achieving $98.94 \pm 0.1\%$ on Citeseer and $96.81 \pm 0.6\%$ on Cora. In ‘Neural Link Prediction with Walk Pooling,’ L. Pan et al proposed a GNN architecture that adds attention mechanisms to GNNs, achieving an accuracy of $89.59 \pm 1.58\%$ [3]. These past studies show the efficacy of GNNs on academic publications for node classification and link prediction, and, since these accuracies were certainly impressive, we hope to use these networks as a baseline to examine the key features of different datasets and extract information to create a similar dataset about the subject of machine learning hardware.

III. APPROACH

Two preconstructed datasets, Cora and CiteSeer, are large datasets of scientific machine learning papers. Each dataset contains papers from before 2008, with citations to other papers in the database as well as the vocabulary used. The nodes of the graph are the papers themselves, with topics and vocabulary as individual features, while the directed links

are citations to other papers. Cora classifies each publication into one of seven categories: Theory, Reinforcement Learning, Genetic Algorithms, Neural Networks, Probabilistic Methods, Case Based, and Rule Learning. On the other hand, Citeseer classifies each publication into one of six categories: Agents, Artificial Intelligence, Database, Human Computer Interaction, Machine Learning, and Information Retrieval.

Meanwhile, PubMed is a database of diabetes publications that is classified into one of three classes [4]. As such, this dataset is similar to both Cora and CiteSeer, though it further explores the efficacy of GNNs on datasets outside of machine learning. This performance of this dataset can further validate parameters potentially affecting accuracies, such as edge count, classes, or path length, in addition to explaining differences in different fields of publications.

Web of Science (WoS) is an online database of publication citations and information [5]. This online database created by Clarivate can be queried by genre and shows identifying information such as titles, references, and specific citations. As such, we can utilize the WoS dataset to address the goal of exploring additional datasets in machine learning hardware. Specifically, by downloading papers by (exclusive) topic, we can create a dataset representative of the field of machine learning and compare the performance of our dataset against GNNs alongside other common datasets to analyze its efficacy on these models and its viability as a dataset in an academic setting. Using this popular database, we query their dataset of publications for four categories: Machine Learning on the Edge, Neuromorphic Computing, Spiking Neural Networks, and Machine Learning Accelerator. In using the Web of Science database, we assume that the papers are representative of their assigned class and distinct from related classes. In this regard, the dataset that we curate is unique from other citation networks in that it will hold distinct characteristics and be of machine learning hardware, providing insight into the structure and nature of the machine learning hardware citation network.

Beginning with the provided models like Deep Graph Library, we take inspiration from existing methods, such as preprocessing methods discussed by the Semantic Scholar Academic Graph [6]. By constructing our datasets of publications, we will examine model characteristics of varying types of publication networks. We also explored other datasets to improve the diversity of our datasets. The Microsoft Open Academic Graph [7], a large academic graph dataset, shows a promising example of dataset scale, though it has differences in features, scaling, and overlapping that can interfere with providing information about any one academic field. Here, the data set uses a variety of data types for the nodes, like publications, authors, and venues, with the links between each denoting a correlation.

Following the specified approach to our project goals, the main tasks of our project are as follows. First, we will apply the necessary pre-processing techniques to each dataset for model compatibility. Then, we will create a GNN model using each dataset that accurately predicts paper topics and citations, commenting on the accuracy and effectiveness of GNNs in

node classification or link prediction in academic contexts. For node classification, we will create a GNN model based on the vocabulary and neighbors as embeddings of each node, experimenting with different configurations to find the best model. Finally, we will take the results in terms of graph characteristics and model performance to analyze our dataset as an academic tool. In this regard, our technique is rather scalable, as these steps can be applied to academic datasets of nearly any field and in datasets beyond academic publications as well. However, this is also given the assumption that different academic fields follow similar citation networks to those listed, that these fields can be described by the vocabularies that they exhibit, and that these networks produce a graph that is compatible with neural network models.

IV. EXPERIMENTAL SETUP AND RESULTS

A. Datasets

Since the datasets Cora and Citeseer are included within some deep learning packages, we directly use those for constructing GNNs. Utilizing resources and configurations of included models within StellarGraph, Pytorch Geometric, and Keras, we can also construct predictive models for node classification and link prediction.

In creating our dataset, we query Web of Science for topics in machine learning hardware published between 1970 and 2022. For ease of use, we removed nodes without a DOI and duplicated values between categories so each node only belongs to one category. The collected data includes title, publishing year, abstract, and citations.

While Web of Science shows a close-to-linear growth of publications during this time, there has been an exponential growth of papers in machine learning hardware over the same period. We see a large spike in interest, following machine learning in general after OpenAI launched GPT.

With the text files generated, we constructed a NetworkX graph using DOI values to match paper citations. We also utilized the categories as labels, and the overall graph topography is in Figure A. This graph serves as our dataset once converted into StellarGraph or NetworkX Pytorch Geometric dataset instances as needed.

In total, our graph had 4951 nodes, though only 2605 nodes have non-zero edges. Looking at the in-degrees of each node, we see evidence of a scale-free network, where few nodes are popularly cited, but many are not referenced often. This may also be due to the recentness of these publications, as future papers will likely add on to those numbers. The number of paper citations is also indicative of its impact, with larger in-degrees meaning more outreach and influence.

Removing nodes without edges allows us to focus on the variety in edges and prevents the model from predicting 0 edges each time for high accuracy, the resulting is still disconnected, with an average degree of 5.0672, an average clustering coefficient of 0.1353, and a maximum diameter across all connected subgraphs of 16, with the average undirected path length being 4.6073 edges.

However, from observation, we notice that our graph is much smaller and less connected than the given datasets. As in Figure 1 below, there are lots of unconnected nodes. And, after removing non-zero edges, there are still plenty of nodes that have only one edge to another paper, likely due to their remaining citations not having an associated word vector.

To combat this, we increase the size of the WoS database, querying additional nodes from the WoS database. Then, we take the largest connected component of the larger graph to increase the overall connectivity of the graph.



Fig. 1. Large WoS network, Blue=ML Accelerators, Purple=Edge ML, Green=Neuromorphic Computing, Orange=Spiking Neural Network.

After applying these operations, the total number of nodes increases from 2,605 to 9,605 after removing unusable nodes. This graph also has 52,945 edges and 6,821 features, from applying the same methods of keyword extraction. When increasing the size of the dataset, we were also met with a class imbalance - there were far fewer papers in ML Accelerators in the WoS database, less than two thousand papers, while the other categories had over five thousand papers each. In this regard, we can infer that ML accelerators are less popular (or at least less explored) than other papers. However, since WoS is our only source of data, this could also result from a lack of these papers in the Web of Science database. Since class imbalances can also impact performance on ML models, this also acts as a limitation on the GNN models in use.

Recognizing that the number of features is nearly double that of any of the given datasets, this can also lead to excessive computation time and increased bias in the results from features that supply noise in the machine learning models. So, we then apply principal component analysis (PCA) to these features to select the features that best describe the classifications and links between each node. Here, we apply PCA with an explained variance of 80% to provide a sufficient level of specificity without overfitting the data, reducing the number of features to 760.

Following this, we also update the characteristics of our graph, producing an average degree of 11.09, an average clustering coefficient of 0.158, and a diameter of 18. From the differences in these datasets, we see that the average degree of our dataset is much higher than the given datasets.

However, the clustering coefficient and diameter of the dataset are comparable to those of other datasets. These metrics are compared to the other datasets in Table 1 below.

From this table, we can see differences among the datasets explored in this paper, revealing the structural properties of these different graphs. Primarily, there are many more edges in the updated WoS compared to the other datasets, resulting in a much higher average degree for the updated dataset. However, the clustering coefficient for our dataset is still comparable to other datasets, implying that the field of machine learning hardware tends to have a significant number of citations, but a similar number of these papers forming a well-defined cluster. While Cora and CiteSeer are most similar to the WoS database in terms of content, both relating to machine learning, the graph characteristics of WoS are most similar to PubMed, with a higher degree and lower clustering coefficient. However, another limitation of utilizing WoS as our primary database is that it is unclear whether these characteristics arise as a quality of the database or as a characteristic of the field of ML hardware in general.

TABLE I
CITATION NETWORK CHARACTERISTICS

	<i>Nodes</i>	<i>Edges</i>	<i>Features</i>	<i>Degree</i>	<i>CC^a</i>	<i>MD^b</i>
<i>Cora</i>	2708	5429	1433	3.89807	0.24067	19
<i>CiteSeer</i>	3312	4660	3704	2.81400	0.14255	28
<i>Pubmed</i>	19717	44348	500	4.49632	0.06017	18
<i>Small WoS</i>	2605	6600	4134	5.06717	0.13529	16
<i>Large WoS</i>	2605	6600	4134	5.06717	0.13529	16

^aClustering Coefficient. ^bMaximum Diameter.

To see the distribution of the demographic of the papers in terms of age, we look at the number of papers in our updated WoS database, as in Figure 2 below. From this graph, we see that papers have been released more frequently in recent years. This implies that machine learning hardware is a rising field, which follows given the recent popularity of machine learning as a field in general.




Fig. 2. Growth of Papers Published in ML Hardware over Time.

After applying these changes, we define a t-SNE visualization on the dataset, as pictured in Figure 3. The plot helps to establish underlying patterns and relationships in the dataset on a 2-D plane for reader understandability. From this plot,

we see that there are some strong clusters of papers and some much weaker, though still characteristic.

Then, after finding the edge list of the dataset, we can plot the degree distribution chart of the network, plotting the number of times a paper was cited against the frequency of that number of citations, shown in Figure 5 below.

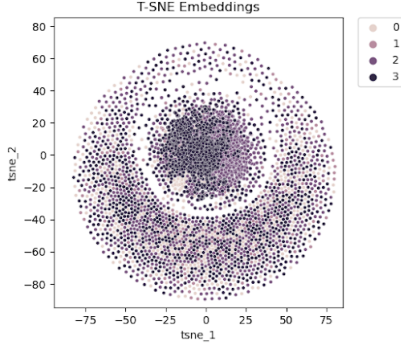


Fig. 3. t-SNE Plot of the WoS Dataset.

From this plot, we establish that the WoS citation network shows characteristics of a small-world network, with few nodes exhibiting a high number of citations and most having very few citations. This result makes sense, as few papers are likely to have a high impact, and those that do are likely to influence many papers.




Fig. 4. Occurrences for papers with certain times cited.

Looking at the most-cited nodes, we find impactful key publications in their respective fields. One such publication is “Equivalent-accuracy accelerated neural-network training using analogue memory”, published in 2016 [7]. In “Nanoscale memristor device as synapse in neuromorphic systems” published in 2018, Jo et al. experimentally show a nanoscale memristor working within the context of neuromorphic computing [8]. In another example, “Memory devices and applications for in-memory computing” gives an overview of potential solutions to in-memory computing, including solutions closely related to machine learning and digital signal processing [9].

In finding additional outliers, we were also able to distinguish the age of our papers. The dataset consisted of papers from the years of 1971 to 2023. Since machine learning was identified as a rising field from the discussion regarding Figure 2, there were many papers from 2023. And, since

many papers only provided a month as the date of publication, the multiple papers from November of 2023 prevented the identification of the newest paper. However, the oldest paper, the only paper from 1971, was titled “Unit Spike Activity in Coelenteran Neural Network”, a paper about temperature spikes in primitive neural nerve networks [10].

While both Cora and Citeseer categorize publications and have vocab dictionaries, there are no such categories or dictionaries for Web of Science results. Therefore, we utilized the search function to specify the categories to be both used and excluded in each search. To match the vocabularies of Cora and Citeseer, we utilized spacy for trained keyword extraction models, finding the top 10 unigrams with the most significance in each abstract, and used these words to construct our vocabulary vectors [11].

Before you begin to format your paper, first write and save the content as a separate text file. Complete all content and organizational editing before formatting. Please note sections IV-B–IV-F below for more information on proofreading, spelling and grammar.

Keep your text and graphic files separate until after the text has been formatted and styled. Do not number text heads— \LaTeX will do that for you.

B. Abbreviations and Acronyms

Define abbreviations and acronyms the first time they are used in the text, even after they have been defined in the abstract. Abbreviations such as IEEE, SI, MKS, CGS, ac, dc, and rms do not have to be defined. Do not use abbreviations in the title or heads unless they are unavoidable.

C. Units

- Use either SI (MKS) or CGS as primary units. (SI units are encouraged.) English units may be used as secondary units (in parentheses). An exception would be the use of English units as identifiers in trade, such as “3.5-inch disk drive”.
- Avoid combining SI and CGS units, such as current in amperes and magnetic field in oersteds. This often leads to confusion because equations do not balance dimensionally. If you must use mixed units, clearly state the units for each quantity that you use in an equation.
- Do not mix complete spellings and abbreviations of units: “Wb/m²” or “webers per square meter”, not “webers/m²”. Spell out units when they appear in text: “. . . a few henries”, not “. . . a few H”.
- Use a zero before decimal points: “0.25”, not “.25”. Use “cm³”, not “cc”.)

D. Equations

Number equations consecutively. To make your equations more compact, you may use the solidus (/), the exp function, or appropriate exponents. Italicize Roman symbols for quantities and variables, but not Greek symbols. Use a long dash rather than a hyphen for a minus sign. Punctuate

equations with commas or periods when they are part of a sentence, as in:

$$a + b = \gamma \quad (1)$$

Be sure that the symbols in your equation have been defined before or immediately following the equation. Use “(1)”, not “Eq. (1)” or “equation (1)”, except at the beginning of a sentence: “Equation (1) is . . .”

E. *L^AT_EX-Specific Advice*

Please use “soft” (e.g., `\eqref{Eq}`) cross references instead of “hard” references (e.g., (1)). That will make it possible to combine sections, add equations, or change the order of figures or citations without having to go through the file line by line.

Please don’t use the `{eqnarray}` equation environment. Use `{align}` or `{IEEEeqnarray}` instead. The `{eqnarray}` environment leaves unsightly spaces around relation symbols.

Please note that the `{subequations}` environment in *L^AT_EX* will increment the main equation counter even when there are no equation numbers displayed. If you forget that, you might write an article in which the equation numbers skip from (17) to (20), causing the copy editors to wonder if you’ve discovered a new method of counting.

BIB_TE_X does not work by magic. It doesn’t get the bibliographic data from thin air but from .bib files. If you use *BIB_TE_X* to produce a bibliography you must send the .bib files.

L^AT_EX can’t read your mind. If you assign the same label to a subsection and a table, you might find that Table I has been cross referenced as Table IV-B3.

L^AT_EX does not have precognitive abilities. If you put a `\label` command before the command that updates the counter it’s supposed to be using, the label will pick up the last counter to be cross referenced instead. In particular, a `\label` command should not go before the caption of a figure or a table.

Do not use `\nonumber` inside the `{array}` environment. It will not stop equation numbers inside `{array}` (there won’t be any anyway) and it might stop a wanted equation number in the surrounding equation.

F. *Some Common Mistakes*

- The word “data” is plural, not singular.
- The subscript for the permeability of vacuum μ_0 , and other common scientific constants, is zero with subscript formatting, not a lowercase letter “o”.
- In American English, commas, semicolons, periods, question and exclamation marks are located within quotation marks only when a complete thought or name is cited, such as a title or full quotation. When quotation marks are used, instead of a bold or italic typeface, to highlight a word or phrase, punctuation should appear outside of the quotation marks. A parenthetical phrase or statement at the end of a sentence is punctuated outside of the closing parenthesis (like this). (A parenthetical sentence is punctuated within the parentheses.)

- A graph within a graph is an “inset”, not an “insert”. The word alternatively is preferred to the word “alternately” (unless you really mean something that alternates).
- Do not use the word “essentially” to mean “approximately” or “effectively”.
- In your paper title, if the words “that uses” can accurately replace the word “using”, capitalize the “u”; if not, keep using lower-cased.
- Be aware of the different meanings of the homophones “affect” and “effect”, “complement” and “compliment”, “discreet” and “discrete”, “principal” and “principle”.
- Do not confuse “imply” and “infer”.
- The prefix “non” is not a word; it should be joined to the word it modifies, usually without a hyphen.
- There is no period after the “et” in the Latin abbreviation “et al.”.
- The abbreviation “i.e.” means “that is”, and the abbreviation “e.g.” means “for example”.

An excellent style manual for science writers is [7].

G. *Authors and Affiliations*

The class file is designed for, but not limited to, six authors. A minimum of one author is required for all conference articles. Author names should be listed starting from left to right and then moving down to the next line. This is the author sequence that will be used in future citations and by indexing services. Names should not be listed in columns nor group by affiliation. Please keep your affiliations as succinct as possible (for example, do not differentiate among departments of the same organization).

H. *Identify the Headings*

Headings, or heads, are organizational devices that guide the reader through your paper. There are two types: component heads and text heads.

Component heads identify the different components of your paper and are not topically subordinate to each other. Examples include Acknowledgments and References and, for these, the correct style to use is “Heading 5”. Use “figure caption” for your Figure captions, and “table head” for your table title. Run-in heads, such as “Abstract”, will require you to apply a style (in this case, italic) in addition to the style provided by the drop down menu to differentiate the head from the text.

Text heads organize the topics on a relational, hierarchical basis. For example, the paper title is the primary text head because all subsequent material relates and elaborates on this one topic. If there are two or more sub-topics, the next level head (uppercase Roman numerals) should be used and, conversely, if there are not at least two sub-topics, then no subheads should be introduced.

I. *Figures and Tables*

a) *Positioning Figures and Tables:* Place figures and tables at the top and bottom of columns. Avoid placing them in the middle of columns. Large figures and tables may span

across both columns. Figure captions should be below the figures; table heads should appear above the tables. Insert figures and tables after they are cited in the text. Use the abbreviation “Fig. 5”, even at the beginning of a sentence.

TABLE II
TABLE TYPE STYLES

Table Head	Table Column Head		
	<i>Table column subhead</i>	<i>Subhead</i>	<i>Subhead</i>
copy	More table copy ^a		

^aSample of a Table footnote.



Fig. 5. Example of a figure caption.

Figure Labels: Use 8 point Times New Roman for Figure labels. Use words rather than symbols or abbreviations when writing Figure axis labels to avoid confusing the reader. As an example, write the quantity “Magnetization”, or “Magnetization, M”, not just “M”. If including units in the label, present them within parentheses. Do not label axes only with units. In the example, write “Magnetization (A/m)” or “Magnetization {A[m(1)]}”, not just “A/m”. Do not label axes with a ratio of quantities and units. For example, write “Temperature (K)”, not “Temperature/K”.

ACKNOWLEDGMENT

The preferred spelling of the word “acknowledgment” in America is without an “e” after the “g”. Avoid the stilted expression “one of us (R. B. G.) thanks ...”. Instead, try “R. B. G. thanks...”. Put sponsor acknowledgments in the unnumbered footnote on the first page.

REFERENCES

Please number citations consecutively within brackets [1]. The sentence punctuation follows the bracket [2]. Refer simply to the reference number, as in [3]—do not use “Ref. [3]” or “reference [3]” except at the beginning of a sentence: “Reference [3] was the first ...”

Number footnotes separately in superscripts. Place the actual footnote at the bottom of the column in which it was cited. Do not put footnotes in the abstract or reference list. Use letters for table footnotes.

Unless there are six authors or more give all authors’ names; do not use “et al.”. Papers that have not been published, even if they have been submitted for publication, should be cited as “unpublished” [4]. Papers that have been accepted for publication should be cited as “in press” [5]. Capitalize only the first word in a paper title, except for proper nouns and element symbols.

For papers published in translation journals, please give the English citation first, followed by the original foreign-language citation [6].

REFERENCES

- [1] G. Eason, B. Noble, and I. N. Sneddon, “On certain integrals of Lipschitz-Hankel type involving products of Bessel functions,” Phil. Trans. Roy. Soc. London, vol. A247, pp. 529–551, April 1955.
- [2] J. Clerk Maxwell, A Treatise on Electricity and Magnetism, 3rd ed., vol. 2. Oxford: Clarendon, 1892, pp.68–73.
- [3] I. S. Jacobs and C. P. Bean, “Fine particles, thin films and exchange anisotropy,” in Magnetism, vol. III, G. T. Rado and H. Suhl, Eds. New York: Academic, 1963, pp. 271–350.
- [4] K. Elissa, “Title of paper if known,” unpublished.
- [5] R. Nicole, “Title of paper with only first word capitalized,” J. Name Stand. Abbrev., in press.
- [6] Y. Yorozu, M. Hirano, K. Oka, and Y. Tagawa, “Electron spectroscopy studies on magneto-optical media and plastic substrate interface,” IEEE Transl. J. Magn. Japan, vol. 2, pp. 740–741, August 1987 [Digests 9th Annual Conf. Magnetism Japan, p. 301, 1982].
- [7] M. Young, The Technical Writer’s Handbook. Mill Valley, CA: University Science, 1989.

IEEE conference templates contain guidance text for composing and formatting conference papers. Please ensure that all template text is removed from your conference paper prior to submission to the conference. Failure to remove the template text from your paper may result in your paper not being published.